

Bayesplot :: CHEAT SHEET




```
library("bayesplot")
library("rstanarm")
library("ggplot2")
library("rstan")
library("dplyr")
```

rstanarm
To showcase bayesplot, we'll fit linear regression using `rstanarm::stan_glm` and use this model throughout.

```
model <- stan_glm(mpg ~ .,
                  data=mtcars, chains=4)
posterior <- as.matrix(model)
```


Chances are good you're most interested in the posterior distributions for select parameters.

```
plot_title <- ggtitle("Posterior distributions",
                      "medians and 80% intervals")
mcmc_areas(posterior,
           pars = c("drat", "am", "wt"),
           prob = 0.8) + plot_title
```




Diagnosing convergence with traceplots is simple.

```
mcmc_trace(posterior, pars=c("sigma", "wt"))
```




The pairs plot is helpful in determining if you have any highly correlated parameters, but necessitates us extracting the stanfit object from our stanreg model.

```
posterior_chains <- rstan::extract(model$stanfit,
                                   inc_warmup = TRUE,
                                   permuted = FALSE)
pairs <- posterior_chains %>%
  mcmc_pairs(pars = c("(Intercept)", "wt", "sigma"))
```




Check how well the model covers your data with draws from the posterior predictive density.

```
model %>%
  posterior_predict(draws = 50) %>%
  ppc_dens_overlay(y = model$y,
                  yrep = .)
```




```
model %>%
  posterior_predict(draws = 500) %>%
  ppc_stat_grouped(y = mtcars$mpg, group = mtcars$am,
                  stat = "median", binwidth=0.5)
```




RStan
Of course, bayesplot plays nicely with RStan, too. We'll just pretend like we fit a model with `rstan::sampling` by extracting the stanfit object from our stanreg model. This allows you to visualize individual chains.

```
fargs <- list(ncol = 2,
              labeller = label_parsed)
pars <- c("(Intercept)", "sigma")
mcmc_trace(posterior_chains, pars = pars,
           n_warmup = 300,
           facet_args = fargs)
```



Diagnostics
Bayesplot makes it easy to check diagnostics specific to the NUTS sampling method that rstanarm uses by default.

```
mcmc_scatter(
  as.matrix(model$stanfit),
  pars = c("wt", "sigma"),
  np = nuts_params(model$stanfit))
```



```
np <- nuts_params(model$stanfit)
mcmc_nuts_energy(np, binwidth=1) +
  ggtitle("NUTS Energy Diagnostic")
```

